

DESIGNING THE TAX ARCHITECTURE OF A GLOBAL INSURANCE GROUP

A 50-Agent, Graph-Based Framework for Headquarters, Regional Platforms, Reinsurance, Capital, and Substance

A technical paper for senior executives, CFOs, tax directors, general counsel, chief risk officers, and international tax advisers. The study develops a synthetic but institutionally realistic framework for selecting and governing a global insurance operating model in the Pillar Two era.

PAPER PURPOSE

The central question is not where the statutory corporate tax rate is lowest. It is where a global insurer can credibly place authority, capital, risk, talent, and systems so that the resulting institution is simpler, resilient, governable, and efficient after minimum tax, withholding leakage, capital friction, regulatory constraints, substance, operating cost, and implementation risk are recognized.

Abstract

This paper develops a decision architecture for redesigning the tax and operating structure of a synthetic multinational insurance group. The proposed model combines deterministic analysis with a 50-agent multidisciplinary review¹ spanning international tax, Pillar Two, transfer pricing, insurance regulation, solvency, actuarial risk, reinsurance, treasury, accounting, corporate law, operations, data, cyber risk and implementation. Deterministic code constructs entity-level data, jurisdiction attributes, ownership and reinsurance networks, candidate headquarters and regional-platform structures, feasibility gates and comparative metrics. Specialist agents then evaluate the same controlled evidence under distinct professional mandates, while structured outputs preserve recommendations, confidence, objections, required evidence and second-order effects.

The framework treats legal compliance, capital adequacy, policyholder protection and genuine management substance as hard constraints rather than variables that can be offset by a low tax rate. It evaluates total structural friction, including minimum-tax effects, withholding leakage, capital immobility, regulatory credibility, operational cost, legal stability, implementation risk and governance complexity. The resulting target model combines one group headquarters with three principal regional platforms while retaining local regulated carriers where licensing, solvency or policyholder obligations require them.

The paper concludes that durable insurance tax planning must align taxable profit with actual authority, financial capacity, risk control, personnel and systems. AI expands issue coverage and structures disagreement, but deterministic validation and accountable human authorization retain final authority. The output is an architecture contract specifying the functions, capital, people, controls and conditions precedent required for implementation.

¹Companion notebook: <https://drive.google.com/file/d/1jNPwx5FKIsYrhqnrqQYmVGSBGZ2QwoZ6/view?usp=sharing>.

The Strategic Problem: Tax Architecture as Institutional Design

A multinational insurer is a uniquely difficult object to restructure. Its balance sheet is not merely a collection of financial assets and liabilities; it is a web of promises to policyholders, regulated capital pools, reserves, reinsurance recoverables, collateral arrangements, investment portfolios and long-duration obligations. Legal entities are often embedded in licensing regimes that cannot be reorganized as freely as ordinary commercial subsidiaries. Moreover, supervisory authorities are concerned not only with solvency at a consolidated level but with the ability of individual regulated entities to meet local obligations under stress.

The board's design problem can therefore be written as a constrained multi-objective optimization:

$$\min_{A \in \mathcal{F}} \Phi(A) = w_T T(A) + w_W W(A) + w_C C(A) + w_O O(A) + w_I I(A) - w_R R(A) - w_S S(A) - w_L L(A), \quad (1)$$

where A is a candidate architecture; T is minimum-tax-adjusted tax friction; W is withholding and cross-border leakage; C is capital friction; O is operating cost; I is implementation complexity; R is regulatory quality; S is substance capacity; and L is legal and institutional stability. The feasible set $\mathcal{F}$ is defined by hard legal, regulatory, solvency, governance and substance constraints.

The distinction between objectives and constraints is fundamental. Tax efficiency is an objective. Compliance with law is a constraint. Capital adequacy is a constraint. Genuine management substance is a constraint. Policyholder protection is a constraint. A scoring model must not allow a sufficiently low tax rate to compensate mathematically for a structure that is legally or economically non-credible.

This design problem is path-dependent. Insurance groups rarely begin with a clean institutional blueprint; they inherit acquisitions, legacy licenses, closed books, historic reinsurance arrangements, tax attributes, minority interests, technology platforms and regulatory commitments created over decades. Some entities continue to exist because they issue new business, while others remain because they administer old portfolios, preserve licenses, hold investments or support guarantees. A credible redesign must therefore distinguish between complexity that is economically necessary and complexity that survives only because no one has tested whether it can be removed. The baseline is not a blank sheet. It is a network of rights and obligations whose unwinding may itself trigger tax, accounting, contractual and supervisory consequences.

The board should consequently require a constraint hierarchy before it considers comparative scores. The first layer contains prohibitions and non-negotiable duties: legality, solvency, licensing, sanctions compliance, policyholder protection and the integrity of financial reporting. The second layer contains feasibility conditions such as regulator consent, treaty entitlement, operational readiness, data localization, employee availability and the financial capacity to assume risk. The third layer contains strategic preferences, including simplicity, capital mobility, cost, speed and tax efficiency. This hierarchy clarifies escalation. A failure in the first layer eliminates a proposal. A weakness in the second layer creates a condition precedent or demands redesign. Only the third layer should be resolved through ordinary trade-offs and weighted comparison.

The analytical boundary must also extend beyond controlled legal entities. Material outsourced administrators, cloud platforms, investment managers, brokers, coverholders, third-party reinsurers, collateral providers and shared-service vendors may be critical to the viability of a proposed structure. A regional platform that looks self-sufficient on an ownership chart may remain dependent on a local carrier's people, data or systems. Conversely, a legally fragmented group may already operate through a small number of centralized capabilities. Mapping these dependencies prevents management from confusing legal simplification with operational simplification. It also reveals where contracts, service-level commitments, resilience testing and exit plans must accompany changes in ownership or functional responsibility.

Time is another design variable. A structure that is feasible in a steady state may be unsafe during transition. For example, a portfolio transfer can require lengthy approval, employee consultation and data migration, while tax-effective dates and accounting cutovers may occur on different timetables. Parallel operating models may be needed for several reporting periods. The group must model not only the destination but also the intermediate states through which it will pass. Each intermediate state should have clear ownership, solvency, liquidity, controls and decision rights. Transition risk belongs inside the objective function because a theoretically superior end state can destroy value if the route to it is excessively fragile, costly or dependent on synchronized approvals that cannot realistically be obtained.

Finally, the optimization should be understood as a decision-support discipline rather than a claim of mathematical certainty. Scores summarize assumptions; they do not eliminate judgment. Sensitivity analysis should reveal whether the ranking changes when tax rates, capital costs, regulatory weights, transition expenses or substance requirements are varied. Scenario analysis should test base, adverse and severe conditions. Decision-makers should see which variables drive the result and which candidates remain attractive across plausible ranges. A robust architecture is not merely the highest-scoring point estimate. It is one whose feasibility and relative advantages persist when reasonable assumptions change and whose principal risks can be governed through explicit commitments, controls and monitoring.

The Synthetic Institution and the Evidence Architecture

Secure analytical environment

The notebook places model credentials in Google Colab Secrets rather than in code and exposes the language-model identifier as a configurable parameter. Live model calls are disabled by default so that the deterministic architecture can be executed without accidental API cost. This separation between computation, credentials and model invocation is a small but important governance pattern for institutional AI.



Figure 1: Notebook Cell 1: environment, reproducibility and secure model configuration.

Why synthetic data is useful

The case study uses synthetic data because the purpose is to develop a reusable decision architecture rather than opine on a named insurer. The synthetic group is sufficiently large and heterogeneous to expose the real conceptual problems: 24 regulated entities, roughly 150 client countries, property and casualty, life, health, specialty, marine, cyber and retirement exposures, plus internal reinsurance links. Premiums, claims, reserves, regulatory capital and employee counts are generated with correlated economic logic rather than as independent random variables.




Figure 2: Notebook Cell 2: synthetic insurer, global footprint and financial profile.


A professional implementation would replace the synthetic table with a controlled data room combining the legal-entity register, regulatory returns, tax provision data, deferred-tax attributes, general ledger, transfer-pricing documentation, actuarial systems, reinsurance administration, treasury data, employment locations, board minutes and decision-rights matrices. The key design principle is *evidence symmetry*: specialists should disagree because their disciplines imply different judgments, not because they were given different facts.

Jurisdictional Decision Surface and the Pillar Two Constraint

The notebook assigns synthetic jurisdictional attributes for corporate tax, withholding friction, treaty quality, regulatory depth, capital fungibility, talent and infrastructure, legal stability, operating cost and the capacity to sustain genuine management substance. The rates are illustrative, not current tax data.

For a group within the scope of the OECD/G20 global minimum-tax framework, a low headline rate cannot be treated as a free-standing advantage. A simplified floor is therefore incorporated into the synthetic scoring model. This is not a GloBE computation. A real model would require jurisdictional effective tax rate calculations based on GloBE income, covered taxes, deferred-tax adjustments, substance-based income exclusion, safe harbours, qualified domestic minimum top-up taxes, IIR and UTPR ordering, ownership structures and local implementation.

The strategic implication is nonetheless clear. As the value of pure rate differentials is compressed, the relative importance of *non-rate architecture* increases: treaty access, capital mobility, regulatory credibility, personnel, infrastructure, legal stability, dispute environment and the ability to demonstrate where key decisions are actually made.







Figure 4: Notebook Cell 4: multilayer ownership and reinsurance graph.

This representation matters because centrality and concentration can reveal dependencies that consolidated accounts obscure. A reinsurance entity may be legally small yet systemically central. A service company may be tax-neutral but operationally indispensable. A local carrier may be capital-rich but unable to remit dividends under stress. The architecture should therefore be optimized over the network, not over a flat table of tax rates.

For senior tax planners, the graph also provides a disciplined way to map transfer-pricing questions. Every material intercompany edge should have an economic characterization: what function is performed, what risk is controlled, what asset is used, what capital is required, what contractual terms exist, and what remuneration follows. The graph becomes a bridge between legal form, economic substance and tax documentation.

The Target Architecture: One Headquarters, Three Regional Platforms

The restructuring target is intentionally ambitious: one listed parent and group headquarters, plus principal platforms for Europe, the Americas and Asia-Pacific. The architecture is not interpreted literally as four legal entities. Local regulated carriers remain beneath the platforms where required. The simplification is therefore a reduction in decision layers, duplicated service structures, unnecessary intercompany relationships and fragmented capital governance.



exposure. A policyholder-protection specialist evaluates whether simplification could weaken local claims-paying capacity or communications. An implementation specialist looks for dependencies, critical paths and sequencing problems. These mandates create purposeful heterogeneity. Without them, fifty calls may simply amplify the same framing assumptions and produce an illusion of consensus.

All specialists work from a common evidence packet. This constraint is essential because it makes disagreement interpretable. If one agent prefers a regional platform and another objects, the divergence should arise from their mandates and professional judgments, not from inconsistent facts. The packet contains the candidate architectures, synthetic jurisdictional attributes, entity-level financial and operational data, ownership and reinsurance graphs, feasibility rules, material assumptions and the definitions of scoring metrics. It also states what is unknown. Missing evidence is not silently completed through speculation; it is converted into an explicit request, caveat or condition precedent. Evidence symmetry thereby supports both auditability and meaningful dissent.

The system separates observation, judgment and authority. Deterministic code observes and constructs the numerical state: it generates the synthetic institution, calculates attributes, builds the graph, enumerates candidates and applies hard gates. Agents exercise bounded judgment over that state. They may interpret trade-offs, identify omitted consequences, question assumptions and recommend further evidence. They do not change the underlying data, waive constraints or directly select the final architecture. Aggregation and validation remain deterministic. This separation mirrors sound organizational governance, in which advisers can influence a decision without possessing unilateral authority to alter capital, contracts or legal obligations.

The response contract is designed to preserve reasoning in a form that can be compared. Each specialist identifies a preferred candidate, states confidence, explains the principal rationale, lists binding constraints, names the most important risk, records an objection to its own recommendation, identifies required evidence, describes a second-order consequence and gives a concise board conclusion. Structured fields support aggregation, while narrative fields retain the context needed for professional review. The schema also makes omissions visible. An output that lacks a candidate identifier, uses invalid values, fails to address the evidence request or omits required coverage is rejected rather than treated as a partial vote.

Confidence is not treated as expertise, probability or truth. It is a self-reported property of an individual judgment and must be interpreted cautiously. A confident answer based on weak evidence should not dominate a cautious answer from a discipline that controls feasibility. The framework therefore combines confidence with domain relevance, evidence quality, vote distribution and red-team findings. It can also cap the influence of any single family of specialists so that numerous closely related tax roles do not overwhelm fewer but decisive regulatory or actuarial roles. Weighting is documented in advance and subjected to sensitivity analysis, reducing the opportunity to tune the process after observing a preferred result.

Independence is similarly treated as a design objective rather than an assumption. Agents may share a base model and therefore remain exposed to correlated reasoning patterns, common blind spots and prompt-induced framing. Mandate diversity mitigates but does not eliminate that dependence. The system should measure similarity in recommendations and rationales, compare results across specialist families and flag apparent consensus that is driven by one repeated argument. Where practical, a professional implementation could introduce model diversity, alternative prompt formulations or independent human reviews for the most consequential issues. The purpose is not to maximize disagreement but to avoid mistaking duplicated reasoning for corroboration.

The institution includes explicit adversarial roles. Some agents are required to argue against the leading candidate, identify how a tax authority or supervisor might characterize it, test whether substance follows profit, and search for dependencies that the scoring model omits. Others examine operational resilience, employee feasibility, data transfers, cyber concentration and political or legal change. These roles are especially important when a candidate performs well numerically. Optimization tends to focus attention on measured variables; red teaming asks what has not been measured, what could invalidate a favorable assumption and how the structure behaves when participants do not accept management's preferred narrative.

The output is best understood as a committee record rather than a referendum. Vote counts provide useful information, but the board should not select a structure merely because it receives the most first preferences. It should examine why agents supported it, which disciplines objected, whether objections share a common cause, and what evidence would resolve the disagreement. A minority objection from group supervision, solvency or policyholder protection may outweigh broad support from roles concerned principally with cost and simplification. The aggregation layer therefore elevates binding constraints and clustered red flags instead of allowing them to disappear inside an average score.

Operationally, the process should run through controlled stages. A preflight confirms model access, schema validity, prompt size and data completeness. A canary call tests the exact production path before the remaining specialists are launched. Responses are stored with model version, prompt version, request identifier, timestamp, latency and validation status. Failed calls are not replaced by invented deterministic opinions. After completion, an exception report identifies missing coverage, invalid outputs and unresolved evidence requests. Only then does aggregation begin. This sequence creates a reproducible chain from evidence to judgment and makes it possible for reviewers to

distinguish analytical conclusions from technical execution failures.

Human professionals remain accountable for disposition. They review the source facts, challenge model interpretations, decide whether objections have been answered and authorize any use of the results. The swarm can expand issue coverage and expose conflicts earlier, but it cannot provide legal opinions, regulatory approvals, actuarial sign-off or fiduciary authorization. The most valuable institutional use is therefore not automatic decision-making. It is the creation of a disciplined, searchable and repeatable forum in which many professional perspectives are applied to the same proposed architecture before management commits capital or begins legal execution.





Figure 6: Notebook Cell 6: the 50-agent multidisciplinary institution.

The design principle is that agent count alone does not create intelligence. Fifty correlated prompts can generate fifty correlated mistakes. Value comes from mandate diversity, common evidence, explicit constraints, structured outputs, adversarial review and deterministic validation. Each agent therefore receives the same evidence pack but a different fiduciary or professional lens.

For a board, this resembles a compressed version of the real governance process. Tax may prefer one architecture, solvency another, operations a third. The system does not suppress disagreement. It records it. The relevant question is not whether the swarm reaches unanimity, but whether the selected architecture can survive the strongest objections of the disciplines whose constraints are binding.

Agent Deliberation and Structured Judgment

The live notebook invokes Claude Sonnet 5 independently for each of the fifty specialists. The production design uses the lowest reasoning-effort setting so that analytical diversity is generated primarily by the fifty distinct professional mandates rather than by unusually long internal deliberation within any single call. Each specialist receives the same evidence pack and returns a schema-constrained judgment containing a preferred architecture, confidence, rationale, hard constraints, principal risk, objection, required evidence, second-order effect and board-level conclusion. The model is expressly instructed to reject designs dependent on evasion, concealment, sham substance, artificial profit allocation, fictitious risk transfer or disregard of insurance regulation.






Figure 8: Notebook Cell 8: weighted consensus, dissent and deterministic red-team penalties.

A useful board-level formulation is to think in terms of a Pareto frontier. An architecture is dominated if another feasible architecture improves at least one material dimension without worsening any other. The board should spend its attention on the non-dominated set, where real strategic trade-offs remain.

The tax function should also distinguish four forms of value:

- **Rate value:** differences in effective taxation after minimum-tax effects.
- **Flow value:** reduced withholding, cash leakage and repatriation friction.
- **Balance-sheet value:** improved use and mobility of regulatory capital, losses and tax attributes.
- **Complexity value:** lower compliance cost, fewer intercompany arrangements, clearer governance and reduced controversy risk.

In mature structures, the last three may be more durable than rate value. They also tend to be more defensible because they arise from genuine simplification and business alignment.

Reinsurance as a Tax, Capital and Risk-Transfer Architecture

Reinsurance deserves separate treatment because it is simultaneously an insurance instrument, a capital-management mechanism, a risk-transfer contract and a potentially material intercompany tax flow. The tax architecture cannot be designed independently of the reinsurance architecture.

The starting point is the commercial purpose of each arrangement. A treaty may protect a local carrier against volatility, provide catastrophe capacity, support entry into a new line, reduce concentration, improve capital efficiency or give access to specialist underwriting and claims expertise. Those purposes determine the economically significant functions and should be visible in pricing, contract terms and actual conduct. A transaction described as risk transfer but operated as a routine funding flow will attract scrutiny from both supervisors and tax authorities. Conversely, a genuine risk-transfer program may still fail if the assuming entity lacks capital, information, experienced personnel or the practical authority to accept, price, monitor and settle the business it has assumed.

The architecture should map the full life cycle of reinsured risk. That map begins with product design and original underwriting, continues through treaty negotiation, pricing, acceptance, premium collection, exposure monitoring, reserving, claims handling, commutations and recoveries, and ends with capital release or final settlement. For each stage, the group should identify who makes decisions, who provides analysis, who can approve exceptions and

who bears the financial consequences. Contractual language is necessary but not sufficient. If a regional reinsurer nominally accepts risk while substantive decisions continue to be made by personnel in the cedant or headquarters, the legal allocation and the operating reality may diverge.

Financial capacity must be assessed in both ordinary and stressed conditions. An assuming entity may appear adequately capitalized under a base forecast yet become dependent on group support after a catastrophe, reserve deterioration, market shock or correlation event. The model should therefore examine solvency capital, liquidity, collateral, retrocession, guarantees, parental support and access to contingent funding. It should also test whether diversification benefits remain recognized by the relevant supervisors and whether local regulators would restrict recoverables, dividends or intra-group exposures during stress. Tax benefits that rely on a reinsurer's ability to absorb risk are not durable if the capital plan assumes that losses will ultimately be returned to the cedants or funded elsewhere in the group.

Pricing analysis must reflect the economics of the treaty rather than a target allocation of profit. Depending on the arrangement, relevant methods may examine expected loss, expenses, cost of capital, risk margins, comparable market pricing, brokerage, collateral and the value of diversification. Long-tail business introduces uncertainty around development patterns and reserve risk; catastrophe business introduces tail concentration and reinstatement features; life and annuity business can involve lapse, longevity, investment and asset-liability risks. A single formula applied across these products may be administratively convenient but economically weak. The transfer-pricing documentation should explain how the method responds to the risk characteristics and why the resulting return is consistent with the functions and financial capacity of the parties.

Withholding, premium taxes, indirect taxes and regulatory charges can materially change treaty economics. The legal characterization of a payment, the residence and licensing status of the reinsurer, local deductions, treaty provisions and beneficial-ownership requirements may all matter. Cash timing also matters because collateral, funds-withheld balances, trust arrangements and delayed recoveries can create financing costs even where the nominal premium rate is attractive. A complete evaluation therefore models the after-tax cash flow of the treaty over time, not merely its impact on a consolidated effective tax rate. It should also recognize the possibility that top-up tax or local minimum taxes reduce the anticipated benefit associated with a low-tax reinsurance location.

Counterparty concentration and operational dependency require explicit limits. Centralizing many treaties in one regional reinsurer can simplify governance and create diversification, but it can also make that entity systemically important to the group. A failure in its systems, claims process, collateral management or regulatory status could affect many local carriers simultaneously. The graph model helps identify this centrality. Management can then establish exposure thresholds, recovery-time objectives, alternative servicing arrangements, retrocession protections and escalation triggers. These controls convert a conceptual platform into a resilient operating capability and provide evidence that centralization was designed with policyholder outcomes in mind.

The same discipline applies to capital fungibility. Regulatory capital is not interchangeable cash, and reinsurance does not automatically make it so. A treaty may reduce required capital in one carrier while increasing collateral or capital needs elsewhere. Recognition may differ across regimes, and supervisors may apply limits, haircuts or approval conditions. The group should compare gross and net solvency positions, local dividend capacity, group diversification, counterparty charges and stress recoverability. It should also model the consequences of ratings pressure, resolution actions or restrictions on intra-group payments. The appropriate objective is improved group resilience and distributable value after these effects, not the mechanical movement of premium to a preferred entity.

Governance should integrate tax, actuarial, risk, finance and regulatory review before treaties are executed or materially amended. Approval papers should state the commercial objective, risk analysis, pricing method, capital effect, tax treatment, collateral terms, accounting consequences, operational responsibilities and performance measures. Material deviations between expected and actual loss experience should trigger review rather than being explained only after a controversy arises. Board and committee minutes should demonstrate that decision-makers understood the downside as well as the projected benefits. This contemporaneous evidence is particularly important for related-party arrangements because it shows that the transaction was evaluated as insurance and capital management, not simply as a vehicle for reallocating taxable income.

Finally, the reinsurance architecture should remain adaptable. Business mix, catastrophe exposure, market capacity, tax rules, accounting standards and supervisory expectations change over time. A treaty that was appropriate when written may become inefficient or risky as portfolios develop. The group should maintain renewal, repricing, commutation and exit strategies and should periodically test whether the location and capabilities of the reinsurer remain suitable. Durable design depends less on fixing one permanent flow than on establishing decision rights and analytical processes that can adjust the program without breaking the alignment among risk, capital, operations and tax.

A regional reinsurance platform may create diversification and capital efficiency, but only if the entity assuming risk has the capital, governance, underwriting expertise, claims oversight and decision-making capacity required to control that risk. The pricing of related-party reinsurance should reflect the actual allocation of economically significant risks and functions. A purely contractual concentration of premium without corresponding control is precisely the

kind of mismatch that sophisticated tax and regulatory authorities examine.

The board should therefore ask five questions for every material internal reinsurance relationship: Who originates the risk? Who prices it? Who decides whether to accept or cede it? Who controls claims and reserving? Who has the capital to absorb adverse outcomes? The answers should align with both the legal contracts and the transfer-pricing model.

Capital Fungibility, Solvency and the Tax Cost of Trapped Capital

Tax models often treat cash as if it were freely movable. Insurance balance sheets make that assumption dangerous. Regulatory capital may be legally owned by the group but economically unavailable for distribution. Dividend restrictions, local solvency requirements, supervisory expectations, collateral requirements and stress conditions can transform nominal cash into trapped capital.

A sophisticated architecture should therefore calculate an after-tax distributable value rather than accounting profit alone. One stylized measure is:

$$V_A = \sum_{t=1}^T \frac{CF_t(A) - Tax_t(A) - Leakage_t(A) - CapitalCost_t(A) - Compliance_t(A)}{(1+r)^t} - TransitionCost(A). \quad (2)$$

This formulation forces the analysis to recognize that a tax saving which requires permanently higher trapped capital may be economically negative.

The capital function and tax function should jointly model dividends, capital injections, reinsurance recoverables, guarantees, debt capacity and stress scenarios. The optimal tax structure is not credible unless it survives the group's capital plan.

Substance, Decision Rights and Transfer Pricing

Substance should not be reduced to headcount. The relevant question is whether the entity has the capacity and authority to perform the functions and control the risks attributed to it. For a regional insurance platform, that may require senior executives, underwriting governance, actuarial capability, finance, treasury, risk management, reinsurance expertise, systems, data and a board that actually exercises authority.

Transfer pricing then becomes a description of the economic institution rather than an ex post allocation exercise. The functional analysis should map people, decision rights, assets, capital and risk control. Routine services can be remunerated as routine services. Strategic risk-taking functions should not be priced as low-value support. Financial transactions and reinsurance require their own analysis.

For senior management, a useful discipline is to require every major profit pool to answer a simple question: *what economically significant activity explains why this profit belongs here?* If the answer cannot be expressed in operational terms that would be credible to a regulator and an auditor, the tax architecture is incomplete.

Red-Team Review: Testing the Architecture Against Failure


The notebook separates proposal generation from criticism. The red-team layer searches for structures whose apparent attractiveness is driven by weak substance, poor capital mobility, inadequate regulatory depth, high withholding friction, legal instability or dependence on a sub-minimum-tax assumption.

A real transaction should extend the red team to controlled foreign company rules, anti-hybrid rules, permanent establishments, beneficial ownership, treaty entitlement, exit taxation, stamp duties, insurance premium taxes, VAT/GST, deferred tax, tax losses, financial assistance, debt covenants, sanctions, data localization, employee transfers, regulatory change-of-control, portfolio-transfer law and policyholder communications.

The red team should also challenge the *narrative*. Can management explain the structure in one coherent story to tax authorities, insurance supervisors, shareholders and employees? A structure requiring a different economic story for each audience is inherently fragile.

The Architecture Contract

The final output should be more than an organization chart. Code Unit 9 filters the evaluated set to architectures that remain hard-feasible and whose deterministic red-team penalty does not exceed the permitted threshold. The best remaining candidate is then translated into an institutional operating model: listed parent and group headquarters, three regional platforms, retained local carriers, reinsurance rules, shared-service principles, substance requirements, tax constraints and conditions precedent. The resulting contract specifies not merely where entities sit, but the institutional conditions under which the selected structure remains credible.



regulatory submissions do not reflect the final operating model. Integrated governance is therefore a value-protection mechanism rather than an administrative overlay.

The program should use decision gates with documented entry and exit criteria. An evidence gate confirms that entity, contract, employee, system, tax and capital inventories are complete and reconciled. A design gate confirms that the target model is legally feasible, financially viable and supported by current analysis. A regulatory gate confirms that required consultations, applications and approvals are progressing on realistic assumptions. An operational-readiness gate tests people, processes, systems, data, controls and contingency arrangements. A closing gate authorizes legal execution only when conditions precedent are satisfied. A stabilization gate confirms that post-close balances, filings, services and governance are operating as intended. These gates prevent schedule pressure from converting unresolved assumptions into operational facts.

Data lineage is especially important because the restructuring will rely on information drawn from systems built for different purposes. Legal-entity registers, policy administration, reinsurance administration, general ledgers, tax provision systems, regulatory returns, actuarial models, treasury records and human-resources data may classify the same entity or transaction differently. The program should define authoritative sources, reconcile material differences and preserve the transformation from source data to board metrics. Changes to assumptions should be version-controlled. This discipline supports Pillar Two calculations, transfer pricing, solvency analysis, accounting entries, regulatory submissions and future audit. It also prevents the AI-assisted analytical layer from appearing more precise than the underlying records justify.

Regulatory engagement should begin before the final structure is presented as irreversible. Supervisors may have concerns about local governance, outsourcing, intra-group reinsurance, dividend capacity, operational resilience, recovery and resolution, policyholder communications or the concentration of decision-making. Early engagement allows those concerns to shape the architecture while alternatives remain available. The group should coordinate messages across jurisdictions without assuming that every supervisor has the same legal mandate or risk appetite. Commitments made during approval processes should be recorded centrally and translated into the operating model, controls and monitoring. An approval obtained through assurances that are not implemented creates both regulatory and governance risk.

Transition-state design deserves the same rigor as the destination. For a period, the group may operate legacy and target processes simultaneously, maintain transitional service agreements, duplicate reporting, retain temporary capital buffers or use interim governance committees. Each temporary arrangement should have an owner, service standard, cost, control framework and termination date. The plan should identify what happens if an approval is delayed, a portfolio transfer is rejected, a key employee does not relocate or a system migration fails. Contingency planning is not evidence of weak commitment; it is recognition that multinational insurance transformations depend on external decisions and complex operational sequences.

The benefits case should be measurable and should distinguish recurring benefits from accounting presentation effects and one-time gains. Expected tax savings, withholding reductions, capital benefits, entity-cost reductions and service efficiencies should be linked to assumptions and owners. Transition taxes, professional fees, systems work, redundancy, recruitment, dual running and stranded costs should be included. Benefits should be risk-adjusted for timing and execution probability. After implementation, actual outcomes should be compared with the approved case, and unexplained variances should trigger review. This prevents the restructuring from being declared successful merely because legal milestones were completed.

Communication must address several audiences with different concerns. Employees need clarity about roles, location, reporting lines and consultation. Regulators need evidence of solvency, governance and policyholder protection. Investors need a credible explanation of costs, benefits, risks and timing. Tax authorities and auditors need consistent documentation of purpose, substance and pricing. Policyholders and distributors may require notice of portfolio transfers, servicing changes or legal counterparties. Although the level of detail will differ, the economic narrative should remain coherent. Contradictory explanations are a warning that the operating model has not been fully resolved.

The roadmap ends with stabilization rather than legal closing. New committees must meet, delegated authorities must operate, intercompany agreements must be followed, service levels must be measured and tax and regulatory filings must reflect the new structure. The group should test whether decision-making has actually moved to the intended locations and whether the relevant personnel possess information and authority. It should also monitor unexpected capital trapping, withholding, tax controversy, operational incidents and employee attrition. A formal post-implementation review should determine whether conditions precedent and regulator commitments remain satisfied and whether corrective actions are required. Only after this period should temporary governance and dual-running arrangements be retired.




Figure 10: Notebook Cell 10: board report, implementation roadmap and provenance export.

Phase I: Evidence and reconciliation

The first phase is not tax design; it is truth-finding. The group should reconcile the legal-entity register, licenses, policy portfolios, reinsurance contracts, tax attributes, deferred tax, capital, guarantees, financing, employees, decision rights, systems and intercompany agreements. Material discrepancies between legal documentation and operating reality should be resolved before optimization.

Phase II: Design validation

Candidate structures should be tested with current jurisdiction-specific tax and legal opinions, a full Pillar Two model, transfer-pricing functional analysis, solvency and actuarial modeling, accounting analysis and supervisory pre-consultation. At this stage, the group should quantify not only recurring tax effects but transition taxes, exit charges, stamp duties, indirect taxes, transaction costs and stranded tax attributes.

Phase III: Build substance before moving profit

Where a regional platform is intended to perform substantive functions, people, systems, governance and decision rights should be built deliberately. The sequencing matters. A credible structure should not depend on attributing profit to a platform before it has the operational capacity to earn that profit.

Phase IV: Execute legal and contractual migration

Only after approvals and operational readiness should the group migrate ownership, contracts, portfolios, reinsurance, financing and service arrangements. Legal steps must be coordinated with regulatory approvals, employee transfers, data migration, accounting cutovers and tax-effective dates.

Phase V: Stabilize and monitor

Post-implementation monitoring should include jurisdictional effective tax rates, top-up tax, solvency, capital mobility, dividend capacity, reinsurance performance, transfer-pricing outcomes, service quality, claims metrics, operational resilience and compliance with commitments made to regulators.

Board Metrics and Decision Dashboard

Senior management should resist the temptation to judge the restructuring by a single tax-savings number. A more useful dashboard includes: jurisdictional and consolidated cash tax; Pillar Two top-up tax; withholding leakage; distributable cash; trapped capital; solvency ratios; return on allocated capital; reinsurance concentration; entity count; number of material intercompany agreements; transfer-pricing controversy exposure; regulatory approvals outstanding; transition cost; implementation milestones; and operational service levels.

A particularly useful metric is the ratio of recurring annual structural benefit to one-time transition cost, adjusted for execution risk. Another is the percentage of group profit for which the location of key decision-makers, risk control and capital can be demonstrated contemporaneously. The latter is not a statutory metric, but it is an excellent internal indicator of structural defensibility.

Governance of the AI-Assisted Process

The use of a 50-agent swarm does not transfer fiduciary responsibility to the model. It changes the analytical process. The system can widen issue coverage, challenge assumptions, expose disciplinary conflicts and accelerate scenario analysis. But material facts must come from controlled sources, calculations must be reproducible, and conclusions must be reviewed by accountable professionals.

The appropriate governance pattern is therefore:

GOVERNANCE SEQUENCE	
Data → Graph → Constraints → Specialists → Debate → Red Team → Deterministic Validation → Human Authorization	

The system should preserve model version, prompt version, source-document version, assumptions, agent outputs, exceptions and human dispositions. In a high-stakes tax project, provenance is not an administrative afterthought; it is part of the evidence that the board followed a disciplined process.

Conclusions

The tax architecture of a multinational insurer is best understood as an institutional operating system. It determines how capital, risk, authority, cash and information move through the group. Tax is embedded in those flows, but it cannot be optimized independently of them.

The move toward a headquarters plus three regional platforms can create substantial value if it reduces duplicated entities, improves capital governance, rationalizes reinsurance, simplifies intercompany arrangements and concentrates genuine expertise. But the architecture must preserve local regulatory legitimacy and align taxable profit with substantive economic functions. In the Pillar Two environment, the strategic premium shifts away from nominal-rate engineering and toward defensible institutional design.

For CEOs and CFOs, the central question is therefore not, “Where is the lowest tax rate?” It is, “Where can this group credibly place authority, capital, risk and talent so that the resulting institution is simpler, more resilient, more governable and more tax-efficient after all constraints are recognized?” For tax directors and advisers, this elevates the function. The highest-value tax planning is no longer a peripheral exercise in entity selection. It is part of the architecture of the enterprise itself.

The enduring output should also include institutional memory. Assumptions, approvals, regulator commitments, implementation decisions and rejected alternatives should be retained in a controlled repository and linked to the architecture contract. Future management teams must be able to understand why functions were located in particular entities and which facts supported that allocation. Without this record, gradual organizational change can separate legal form from operating reality and recreate the same fragmentation the restructuring was intended to remove. Periodic architecture reviews should therefore form part of ordinary governance, especially after acquisitions, major product changes, regulatory reforms, capital events or significant movements of senior personnel.

References

1. OECD, *Global Anti-Base Erosion Model Rules (Pillar Two)*, OECD/G20 Inclusive Framework on BEPS. <https://www.oecd.org/tax/beps/tax-challenges-arising-from-the-digitalisation-of-the-economy-global-anti-base-erosion-model-rules-pillar-two.htm>
2. OECD, *OECD Transfer Pricing Guidelines for Multinational Enterprises and Tax Administrations 2022*. <https://doi.org/10.1787/0e655865-en>
3. OECD, *Transfer Pricing*. <https://www.oecd.org/en/topics/transfer-pricing.html>
4. HM Revenue & Customs, *Multinational Top-up Tax and Domestic Top-up Tax Manual*. <https://www.gov.uk/hmrc-internal-manuals/multinational-top-up-tax-and-domestic-top-up-tax>
5. HM Revenue & Customs, *How to report Pillar 2 Top-up Taxes*. <https://www.gov.uk/guidance/how-to-report-pillar-2-top-up-taxes>
6. International Association of Insurance Supervisors, *Insurance Core Principles and Common Framework for the Supervision of Internationally Active Insurance Groups (ComFrame)*. <https://www.iaisweb.org/activities-topics/standard-setting/icps-and-comframe/>
7. International Association of Insurance Supervisors, *Insurance Capital Standard*. <https://www.iaisweb.org/activities-topics/standard-setting/insurance-capital-standard/>
8. IFRS Foundation, *IFRS 17 Insurance Contracts*. <https://www.ifrs.org/issued-standards/list-of-standards/ifrs-17-insurance-contracts/>
9. Bank of England, Prudential Regulation Authority, *Insurance Supervision*. <https://www.bankofengland.co.uk/prudential-regulation>
10. European Insurance and Occupational Pensions Authority, *Solvency II*. https://www.eiopa.europa.eu/browse/regulation-and-policy/solvency-ii_en
11. National Association of Insurance Commissioners, *International Insurance Relations*. <https://content.naic.org/insurance-topics/international-insurance-relations>
12. Financial Conduct Authority, *Insurance*. <https://www.fca.org.uk/firms/insurance>
13. Anthropic, *API Documentation*. <https://docs.anthropic.com/>

Mapping of the Ten Notebook Cells to the Paper

Unit	Function in the analytical architecture
1	Secure environment, reproducibility and model configuration.
2	Synthetic global insurer and economic footprint.
3	Jurisdiction attributes, minimum-tax adjustment and compliance gates.
4	Multilayer ownership and reinsurance graph.
5	Candidate headquarters and regional architectures.
6	Fifty differentiated professional agents.
7	Fifty live Claude Sonnet 5 specialist calls using schema-constrained outputs, low reasoning effort, preflight/canary validation and no deterministic substitution.
8	Weighted consensus, dissent aggregation and deterministic red-team penalties.
9	Selection of the eligible architecture and translation into an institutional operating-model contract.
10	Board report, five-phase implementation roadmap and provenance/export package.

Professional Validation Checklist

PROFESSIONAL VALIDATION REQUIRED

Before using the framework for an actual group, management should require verified current-law analysis for corporate income tax, Pillar Two, CFC rules, transfer pricing, permanent establishments, beneficial ownership, treaties, withholding taxes, insurance premium taxes, VAT/GST, exit taxes, stamp duties and tax attributes; actuarial and solvency modeling; reinsurance legal analysis; regulatory approvals; securities-law analysis; accounting and deferred-tax treatment; employment and mobility; data protection; sanctions; operational resilience; and a complete implementation plan. The synthetic model is a decision architecture, not a substitute for those workstreams.